

Assembly Line Intelligence

REINFORCE Policy Gradient
Production Decision Optimization

REINFORCE

POLICY GRADIENT

NUMPY · PANDAS

S1

S2

S3

S4

S5

S6

6-station assembly · 11 state features · 6 interventions



23,832

Baseline steps

94.86

Random reward/step

+32.1%


Reward improvement

+20.6%

Throughput gain

The Baseline — What Random Management Looks Like

23,832 decision steps. Every action chosen uniformly at random. This is the performance floor.



Reward σ = 22.97 on mean = 94.86 $\rightarrow$ CoV = 24.2%. Wide variance, zero learning trend across 300 episodes.

94.86

Mean
reward / step

22.97

Reward
std deviation

93.6

Mean
throughput UPH

300



Training episodes

Environment & REINFORCE — Technical Specification

A softmax policy over 11 features. Parameters updated by gradient ascent on expected return.

0 — Reassign Operator

↑ operator efficiency · ↓ micro-stops

1 — Increase Speed

↑ throughput · ↑ failure risk

2 — Decrease Speed

↓ throughput · ↓ failure risk

3 — Quick Maintenance

↓↓ micro-stops · ↓ failure probability

4 — Redirect Flow

↓ WIP accumulation · ↑ line speed

5 — No Action

Baseline — no intervention

Policy: $\pi_{\theta}(a|s) = \text{softmax}(W \cdot s + b)$

$W \in \mathbb{R}^{(6 \times 11)}$ · $b \in \mathbb{R}^6$ · θ initialized to 0
At init: all 6 actions equally likely (16.7%)

Update: $\theta \leftarrow \theta + \alpha \cdot G_t \cdot \nabla \log \pi_{\theta}(a_t|s_t)$

High $G_t \rightarrow$ increase prob of a_t in s_t
Low $G_t \rightarrow$ decrease it

Hyperparameters

gamma (γ)

0.99

Future rewards nearly as valuable as immediate

lr (α)

0.002

Small steps — PG estimates are high-variance

episodes

400

Sufficient for stable convergence

max_steps


100

Episode cap — realistic shift window

REINFORCE operates on full-episode Monte Carlo returns — noisier than TD but unbiased.

400 Episodes — The Gradient Ascending

From 9,873 to 12,727 mean episode reward — +28.9% purely from policy gradient updates.



Total improvement:

9,873



12,727

+28.9% in episode reward · +32.1% in step reward · $P(\text{REINFORCE} > \text{Random}) = 100\%$

Ep 1–50

9,873

Ep 50–150

11,144

Ep 150–250

11,996

Ep 250–400

12,714

Note: High variance is expected in REINFORCE. Monte Carlo returns are noisier than TD estimates — this is a property of the algorithm, not instability.


Variance reduction methods (baseline subtraction, Actor-Critic) are the natural next step from this project.

What the Agent Learned to Do

Learned action probabilities vs. random baseline — across 500 sampled states.



Policy Weight Matrix W — feature influence per action:



Green weights push the agent toward that action. Red weights push it away. The pattern is operationally interpretable.

Redirect Flow

30.1% (+13.4pp)

Dominant action — suppresses WIP buildup before terminal penalty

Reassign Operator

24.3% (+7.6pp)

Second priority — efficiency multiplies directly into throughput

Decrease Speed

4.1% (-12.6pp)

Strongly suppressed — rarely the optimal first response

REINFORCE vs. Random — The Numbers

30 evaluation episodes, same environment, same seeds. Step-level and episode-level comparison.



Step reward
Random

97.43

Step reward
REINFORCE

128.74

Reward
improvement

+32.1%

UPH
improvement

+20.6%

Bootstrap Validation (n=2,000 resamples):

Mean improvement per episode: +3,766 reward units

95% CI: [+2,851, +4,679] · P(REINFORCE > Random): 100.0%

The CI lower bound is positive — the improvement is not noise.

Every single bootstrap resample showed REINFORCE outperforming the random baseline.

Three Operational Scenarios — The Policy in Action

`recommend_action(state)` → top intervention with probability and two alternatives.

A — High WIP Congestion

wip_norm=0.75 · failure_prob=0.05 · operator_efficiency=0.88



1

Redirect Flow

p=0.297

↓ WIP directly before terminal penalty

2

Reassign Operator

p=0.237

Prevent efficiency loss

3

Quick Maintenance

p=0.189

Secondary — address micro-stops

High WIP triggers -1.5 per unit in the reward function. Redirect Flow is the only action that directly reduces WIP. The policy correctly prioritizes it over speed changes.

B — High Failure Risk

failure_prob=0.18 · micro_stops_norm=0.70 · wip_norm=0.20



1

Redirect Flow

p=0.303

Maintains flow stability

2

Reassign Operator

p=0.247

Efficiency support

3

Quick Maintenance

p=0.186

Addresses both failure & micro-stops

failure_prob at 0.18 is near the terminal threshold (0.20). The -10 penalty per unit is the heaviest term. Quick Maintenance reduces both failure_prob and micro_stops simultaneously.

C — Low Operator Efficiency

operator_efficiency=0.68 · ct_norm_s1-s6=0.80 · speed_norm=0.72



1

Redirect Flow

p=0.335

Flow management first

2

Reassign Operator

p=0.258

High-leverage intervention at 0.68 efficiency

3

Quick Maintenance

p=0.173


Secondary

Operator efficiency multiplies directly into throughput: $(3600/ct) \times vel \times efficiency$. At 0.68, Reassign Operator delivers the highest marginal gain. Decrease Speed is suppressed — the policy learned it is not a personnel solution.

Decrease Speed is suppressed in all three scenarios — the policy learned it is rarely the correct first response.

Hyperparameter Sensitivity — Learning Rate

Last-50-episode mean reward across learning rate values (200 training episodes each, seed 42).



Five Conclusions

1. The policy function scales where the Q-table cannot — 66 parameters generalize across all 100+ unseen states.
2. Return normalization is not optional — without it, high-return episodes dominate gradient steps and the policy fails to learn from the rest.
3. The learned priorities are operationally defensible — Redirect Flow (30%) and Reassign Operator (24%) match what an experienced line manager would prioritize.
4. $P(\text{REINFORCE} > \text{Random}) = 100\%$ across 2,000 bootstrap resamples — the improvement is not variance.
5. High reward variance is expected in REINFORCE — Monte Carlo returns are noisier than TD. The response is baseline subtraction or Actor-Critic, not a larger learning rate.

Q-Learning remembered. REINFORCE decided.

The difference is not academic. A Q-table scales poorly with continuous states — each new dimension multiplies the table size. A policy function generalizes gracefully through shared parameters W and b .


REINFORCE trained on this assembly line learned three things a random manager never figures out: High WIP $\rightarrow$ redirect flow. High failure risk $\rightarrow$ maintenance beats throughput. Low efficiency $\rightarrow$ reassign before touching anything else.

One chapter remains: Model-Based RL.

Where Q-Learning reacted and REINFORCE decided, the model-based agent will plan — simulating futures before acting.

Because the best decisions are made before the situation demands them.

+32.1% reward · +20.6% UPH



REINFORCE Policy

400 episodes · $\alpha=0.002$ · $\gamma=0.99$